



deeplearning.ai

Setting up your ML application

Train/dev/test sets

Applied ML is a highly iterative process

layers


hidden units

learning rates

activation functions

...

Idea




Experiment

Code

NLP, Vision, Speech, Structured Data


└─┬─┘
└─┬─┘
Ads Search Security Logistic ...

Train/dev/test sets




Mismatched train/test distribution

Certs



→ Make sure dev and test come from same distribution.



Not having a test set might be okay. (Only dev set.)






deeplearning.ai

Setting up your
ML application

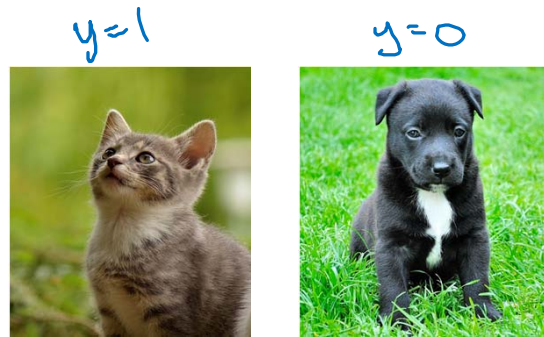
Bias/Variance

Bias and Variance




Bias and Variance

Cat classification



Train set error:	1%	15%	15%	0.5%
Dev set error:	11%	16%	30%	1%
	high variance ↑	high bias ↑↑	high bias & high variance	low bias low variance ↑
Human: ~0%				
Optimal (Bayes) error: ~0%		15%		
		Blurry images		

High bias and high variance





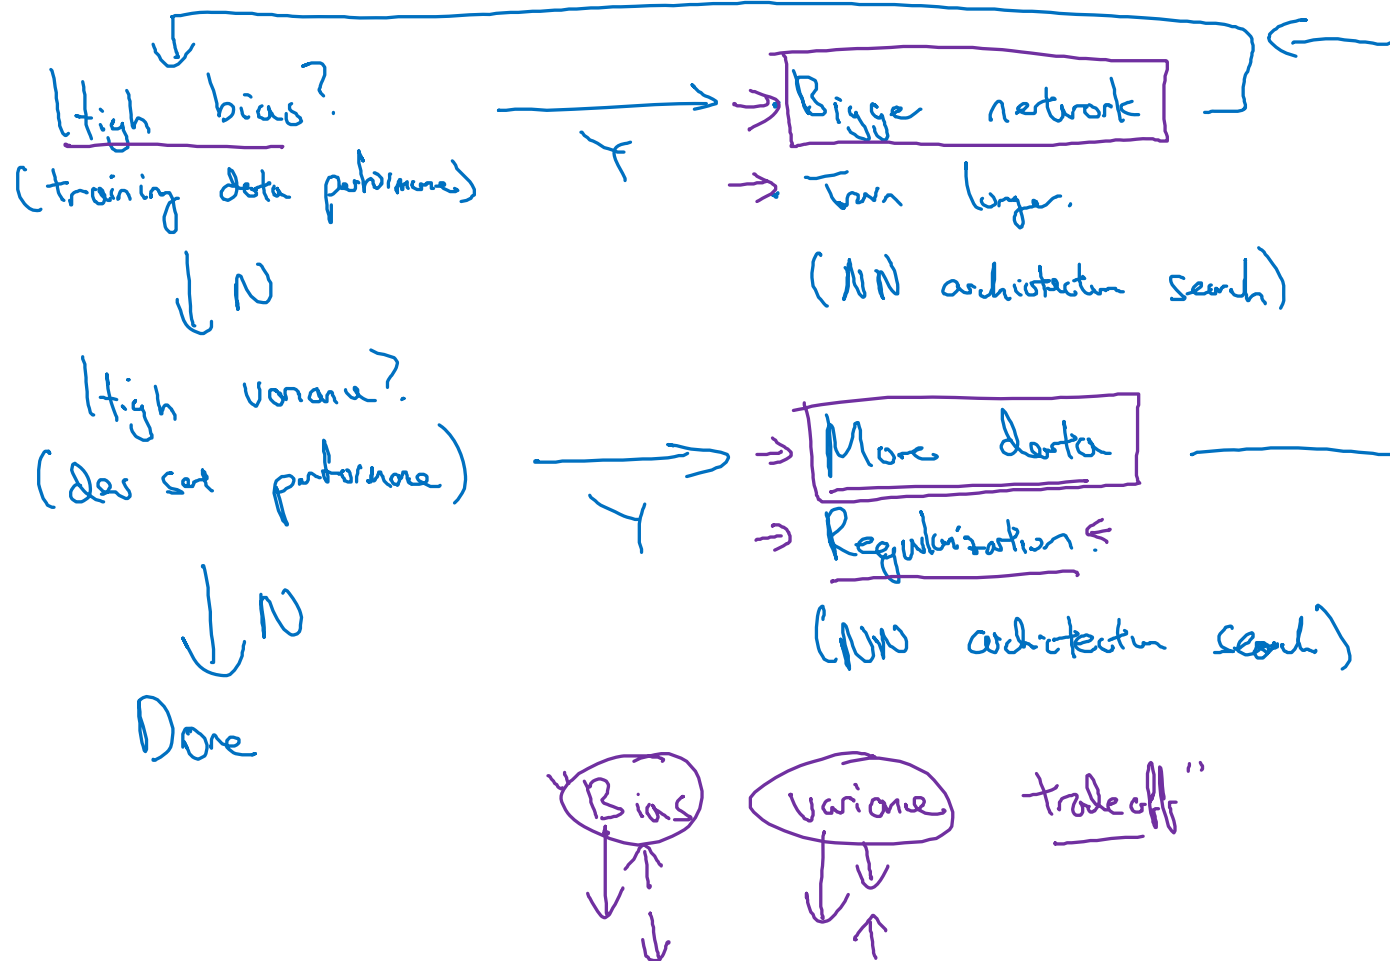
deeplearning.ai

Setting up your
ML application

Basic “recipe”
for machine learning

Basic “recipe” for machine learning

Basic recipe for machine learning





deeplearning.ai

Regularizing your
neural network

Regularization

Logistic regression

$$\min_{w,b} J(w,b)$$

$$\underline{w \in \mathbb{R}^{n_x}}, \underline{b \in \mathbb{R}}$$

$\lambda =$ regularization parameter
lambda lambda

$$J(w,b) = \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(y^{(i)}, \hat{y}^{(i)})}_{\text{cost function}} + \frac{\lambda}{2m} \underbrace{\|w\|_2^2}_{\text{L2 regularization}}$$

~~$$+ \frac{\lambda}{2m} b^2$$~~
 omit

L_2 regularization $\underline{\|w\|_2^2} = \sum_{j=1}^{n_x} w_j^2 = w^T w \leftarrow$

L_1 regularization $\frac{\lambda}{2m} \sum_{j=1}^{n_x} |w_j| = \frac{\lambda}{2m} \|w\|_1$

w will be sparse

Neural network

$$\rightarrow J(w^{[1]}, b^{[1]}, \dots, w^{[L]}, b^{[L]}) = \underbrace{\frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, \hat{y}^{(i)})}_{\text{loss}} + \underbrace{\frac{\lambda}{2n} \sum_{l=1}^L \|w^{[l]}\|_F^2}_{\text{weight decay}}$$

$$\|w^{[l]}\|_F^2 = \sum_{i=1}^{n^{[l]}} \sum_{j=1}^{n^{[l-1]}} (w_{ij}^{[l]})^2$$

$w^{[l]}: \begin{matrix} n^{[l]} & n^{[l-1]} \\ \uparrow & \uparrow \end{matrix}$

"Frobenius norm"

$\|\cdot\|_2^2$

$\|\cdot\|_F^2$

$$dw^{[l]} = \left[\text{(from backprop)} + \frac{\lambda}{n} w^{[l]} \right]$$

$$\rightarrow w^{[l]} := w^{[l]} - \alpha \underbrace{dw^{[l]}}_{\text{weight decay}}$$

$$\underbrace{\frac{\partial J}{\partial w^{[l]}}} = dw^{[l]}$$

"Weight decay"

$$w^{[l]} := w^{[l]} - \alpha \left[\text{(from backprop)} + \frac{\lambda}{n} w^{[l]} \right]$$





$$= w^{[l]} - \frac{\alpha \lambda}{n} w^{[l]} - \alpha \text{(from backprop)}$$

$$= \underbrace{\left(1 - \frac{\alpha \lambda}{n}\right)}_{< 1} \underbrace{w^{[l]}}_{\text{weight decay}} - \alpha \text{(from backprop)}$$

Neural network

$$J(w^{[1]}, b^{[1]}, \dots, w^{[L]}, b^{[L]}) = \frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, \hat{y}^{(i)}) + \frac{\lambda}{2n} \sum_{l=1}^L \|w^{[l]}\|^2$$

How does regularization prevent overfitting?



How does regularization prevent overfitting?




deeplearning.ai

Regularizing your neural network




Why regularization reduces overfitting

How does regularization prevent overfitting?




$$J(w^{(1)}, b^{(1)}) = \frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, \hat{y}^{(i)}) + \frac{\lambda}{2n} \sum_{l=1}^L \|w^{(l)}\|_F^2$$

$$w^{(1)} \approx 0$$



How does regularization prevent overfitting?




$$g(z) = \tanh(z)$$

$\lambda \uparrow$

$W^{[L]} \downarrow$

$$z^{[L]} = W^{[L]} a^{[L-1]} + b^{[L]}$$



Every layer \approx linear.

$$J(\dots) = \underbrace{\sum_i \mathcal{L}(\hat{y}^{(i)}, y^{(i)})}_{\text{loss}} + \underbrace{\frac{\lambda}{2m} \sum_l \|W^{[l]}\|_F^2}_{\text{regularization}}$$





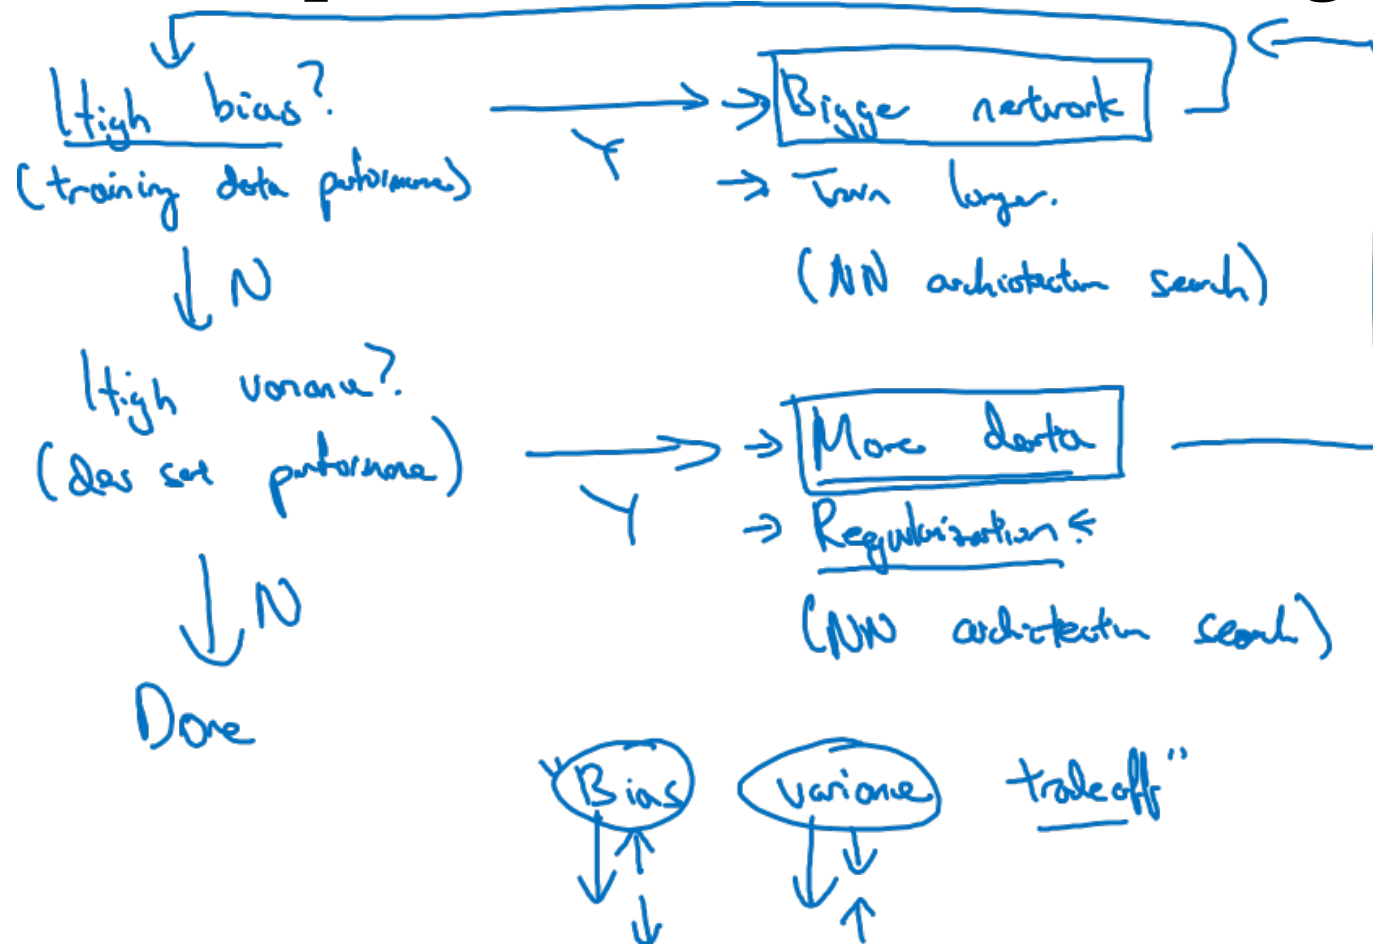
deeplearning.ai

Setting up your ML application

Basic “recipe”
for machine learning

Basic “recipe” for machine learning

Basic recipe for machine learning






deeplearning.ai

Regularizing your neural network

Dropout regularization

Dropout regularization



↑
0.5 ↑
0.5 ↑
0.5

Implementing dropout ("Inverted dropout")

Illustrate with layer $l=3$. $\text{keep-prob} = \frac{0.8}{x}$ 0.2

$\rightarrow \boxed{d3} = \text{np.random.rand}(a3.\text{shape}[0], a3.\text{shape}[1]) < \text{keep-prob}$

$\underline{a3} = \text{np.multiply}(a3, d3)$ $\# a3 \neq d3.$

$\rightarrow \boxed{a3 /= \cancel{0.8} \text{ keep-prob}} \leftarrow$

50 units. \leadsto 10 units shut off

$$z^{[4]} = w^{[4]} \cdot \underline{a^{[3]}} + b^{[4]}$$

\uparrow reduced by 20%.

$$/= \underline{0.8}$$

Test

Making predictions at test time

$$a^{(0)} = X$$


No drop out.

$$z^{(1)} = W^{(1)} \underline{a^{(0)}} + b^{(1)}$$

$$a^{(1)} = g^{(1)}(z^{(1)})$$

$$z^{(2)} = W^{(2)} \underline{a^{(1)}} + b^{(2)}$$

$$a^{(2)} = \dots$$



/= keep-prob





deeplearning.ai

Regularizing your neural network

Understanding dropout

Why does drop-out work?

Intuition: Can't rely on any one feature, so have to spread out weights. \leadsto Shrink weights. b_2






deeplearning.ai


Regularizing your
neural network

Other regularization
methods

Data augmentation



4




4

4

4



Early stopping






deeplearning.ai

Setting up your
optimization problem

Normalizing inputs

Normalizing training sets

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$



Subtract mean:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^m x^{(i)}$$

$$x := x - \mu$$

Normalize variance

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^m x^{(i)} * x^{(i)T}$$

← element-wise

$$x /= \hat{\sigma}^2$$



Use same μ $\hat{\sigma}^2$ to normalize test set.

Why normalize inputs?

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$



Unnormalized:

$w_1: x_1: 1 \dots 1000 \leftarrow$
 $w_2: x_2: 0 \dots 1 \leftarrow$
 $-1 \dots 1$



$x_1: 0 \dots 1$
 $x_2: -1 \dots 1$
 $x_3: 1 \dots 2$

Normalized:






deeplearning.ai

Setting up your
optimization problem

Vanishing/exploding
gradients

Vanishing/exploding gradients



$$g(z) = z, \quad b^{(1)} = 0.$$

$$\hat{y} = w^{(L)} \left(w^{(L-1)} w^{(L-2)} \dots \left(w^{(2)} w^{(1)} x \right) \right) a^{(1)}$$

$$1.5^L$$

$$0.5^L$$

$$w^{(1)} > I$$

$$w^{(2)} < I \quad \begin{bmatrix} 0.9 & \\ & 0.9 \end{bmatrix}$$


$$w^{(2)} = \begin{bmatrix} 0.5 & 0 \\ 0 & 1.5 \end{bmatrix}$$

$$\hat{y} = w^{(L)} \begin{bmatrix} 0.5 & 0 \\ 0 & 1.5 \end{bmatrix}^{L-1} x$$

$$1.5^{L-1} \times$$

$$0.5^{L-1} \times$$

Single neuron example



$$z = w_1 x_1 + w_2 x_2 + \dots + w_n x_n$$

large $n \rightarrow$ Smaller w_i

$$\text{Var}(w_i) = \frac{1}{n} \frac{2}{n}$$

$$\underline{w^{[1]}} = \text{np.random.randn}(\text{shape}) * \text{np.sqrt}\left(\frac{2}{n^{[1-1]}}\right)$$

ReLU $g^{[1]}(z) = \text{ReLU}(z)$

Other variants:

tanh

$$\frac{1}{n^{[1-1]}}$$

Xavier initialization \uparrow

$$\sqrt{\frac{2}{n^{[1-1]} + n^{[1]}}}$$

\uparrow




deeplearning.ai

Setting up your optimization problem

Numerical approximation of gradients

Checking your derivative computation

I $f(\theta) = \theta^3$
 $\theta \in \mathbb{R}.$



$$g(\theta) = \frac{d}{d\theta} f(\theta) = f'(\theta)$$

$\frac{dw}{db}$ \rightarrow $g(\theta) = 3\theta^2$

$g(\theta) = 3 \cdot (1)^2 = 3$
 when $\theta = 1$

$$\frac{f(\theta + \epsilon) - f(\theta)}{\epsilon} \approx g(\theta)$$


$$\frac{(1.01)^3 - 1^3}{0.01} = 3.0301 \approx 3$$

Annotations: 0.0301 is the difference between 3.0301 and 3 . 3.1 and 3.2 are also noted near the result.

$\theta = 1$
 $\theta + \epsilon = 1.01$

Checking your derivative computation

$$\underline{f(\theta) = \theta^3}$$



$$\left[\frac{f(\theta + \epsilon) - f(\theta - \epsilon)}{2\epsilon} \approx \underline{g(\theta)} \right]$$

$$\frac{(1.01)^3 - (0.99)^3}{2(0.01)} = 3.0001 \approx 3$$

$$g(\theta) = 3\theta^2 = 3$$

approx error: 0.0001

(prev slide: 3.0301. error: 0.03)

$$\left\{ \begin{array}{l} f'(\theta) = \lim_{\epsilon \rightarrow 0} \frac{f(\theta + \epsilon) - f(\theta - \epsilon)}{2\epsilon} \quad \begin{array}{l} \mathcal{O}(\epsilon^2) \\ 0.01 \\ \underline{0.0001} \end{array} \end{array} \right. \quad \left| \quad \begin{array}{l} \frac{f(\theta + \epsilon) - f(\theta)}{\epsilon} \quad \text{error: } \mathcal{O}(\epsilon) \\ \uparrow \quad \uparrow \\ \quad 0.01 \end{array} \right.$$



deeplearning.ai

Setting up your
optimization problem

Gradient Checking

Gradient check for a neural network

Take $W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]}$ and reshape into a big vector θ .

concentrate

$$J(w^{[1]}, b^{[1]}, \dots, w^{[L]}, b^{[L]}) = J(\theta)$$

Take $dW^{[1]}, db^{[1]}, \dots, dW^{[L]}, db^{[L]}$ and reshape into a big vector $d\theta$.

concentrate

Is $d\theta$ the gradient of $J(\theta)$?

Gradient checking (Grad check)

$$J(\theta) = J(\theta_1, \theta_2, \theta_3, \dots)$$

for each i :

$$\rightarrow \underline{d\theta_{\text{approx}}[i]} = \frac{J(\theta_1, \theta_2, \dots, \overset{\downarrow}{\theta_i + \epsilon}, \dots) - J(\theta_1, \theta_2, \dots, \overset{\downarrow}{\theta_i - \epsilon}, \dots)}{2\epsilon}$$

$$\approx \underline{d\theta[i]} = \frac{\partial J}{\partial \theta_i} \quad | \quad d\theta_{\text{approx}} \approx d\theta$$

Checks

$$\rightarrow \frac{\|d\theta_{\text{approx}} - d\theta\|_2}{\|d\theta_{\text{approx}}\|_2 + \|d\theta\|_2}$$

$$\underline{\epsilon = 10^{-7}}$$

$$\approx \frac{10^{-7}}{10^{-5}} - \text{great!} \leftarrow$$

$$\rightarrow 10^{-3} - \text{worry.} \leftarrow$$



deeplearning.ai

Setting up your
optimization problem

Gradient Checking
implementation notes

Gradient checking implementation notes

- Don't use in training – only to debug

$$\frac{d\theta_{\text{approx}}[i]}{\uparrow \uparrow} \longleftrightarrow \frac{d\theta[i]}{\uparrow}$$

- If algorithm fails grad check, look at components to try to identify bug.

$$\frac{db_F^{[L]}}{\uparrow} \quad \frac{dW_F^{[L]}}{\uparrow}$$

- Remember regularization.

$$\underline{J(\theta)} = \frac{1}{n} \sum_i \ell(y^{(i)}, \hat{y}^{(i)}) + \underbrace{\frac{\lambda}{2n} \sum_l \|W^{(l)}\|_F^2}_{\text{d}\theta = \text{gradient of } J \text{ w.r.t. } \theta}$$

- Doesn't work with dropout.

$$\underline{J} \quad \underline{\text{keep-prob} = 1.0}$$

- Run at random initialization; perhaps again after some training.

$$\underline{W, b \approx 0}$$